MONASH University

FIT5215 Deep Learning

Quiz for: Convolutional Neural Network

Tutor Team

Department of Data Science and AI
Faculty of Information Technology, Monash University
Email: **trunglm@monash.edu**

GROUP OF EIGHT AUSTRALIA

Question 1

*Given an 3D input tensor with shape [32, 32, 3] over which we apply a conv2D with **16 filters** each of which has shape **[5,5]**, strides **[3,3]**, and padding **valid**. What is the shape of the output tensor?*

- ❑ A. [10, 10]
- ❑ B. [11, 11]
- ❑ C. [11, 11, 16]
- ❑ D. [10, 10, 16]

Question 1

*Given an 3D input tensor with shape [32, 32, 3] over which we apply a conv2D with **16 filters** each of which has shape **[5,5]**, strides **[3,3]**, and padding **valid**. What is the shape of the output tensor?*

- ❑ A. [10, 10]
- ❑ B. [11, 11]
- ❑ C. [11, 11, 16]
- ❑ D. [10, 10, 16] **[x]**

floor operation

$$\left\lfloor \frac{32-5}{3} \right\rfloor + 1 = 10$$

valid means does not use zero padding

Question 2

*Given an 3D input tensor with shape [32, 32, 3] over which we apply a conv2D with **16 filters** each of which has shape **[5,5]**, strides **[3,3]**, and padding **same**. What is the shape of the output tensor?*

- ❑ A. [10, 10]
- ❑ B. [11, 11]
- ❑ C. [11, 11, 16]
- ❑ D. [10, 10, 16]

Question 2

*Given an 3D input tensor with shape [32, 32, 3] over which we apply a conv2D with **16 filters** each of which has shape **[5,5]**, strides **[3,3]**, and padding **same**. What is the shape of the output tensor?*

- ❑ A. [10, 10]
- ❑ B. [11, 11]
- ❑ C. [11, 11, 16] [**x**]
- ❑ D. [10, 10, 16]

$$\left\lfloor \frac{32-1}{3} \right\rfloor + 1 = 11$$

Question 3

*Given an 3D input tensor with shape [64, 64, 10] over which we apply a **max pooling** layer with kernel size **[3,3]**, strides **[3,3]**, and padding **same**. What is the **shape** of the **output tensor**?*

- ❑ A. [21, 21]
- ❑ B. [22, 22]
- ❑ C. [22, 22, 10]
- ❑ D. [22, 22, 3]

Question 3

*Given an 3D input tensor with shape [64, 64, 10] over which we apply a **max pooling** layer with kernel size **[3,3]**, strides **[3,3]**, and padding **same**. What is the **shape** of the **output tensor**?*

- ❑ A. [21, 21]
- ❑ B. [22, 22]
- ❑ C. [22, 22, 10] **[x]**
- ❑ D. [22, 22, 3]

$$\left\lfloor \frac{64-1}{3} \right\rfloor + 1 = 22$$

Question 4

*Assume that the tensor before the last tensor of a CNN has shape [32, 32, 32, 10] and we apply **5 filters** each of which has the shape **[5,5,10]** and strides= **[2,2]** with padding = **'same'** to obtain the last tensor. What is the **shape** of the output tensor?*

- A. [16, 16, 5]
- B. [14, 14, 5]
- C. [32, 14, 14, 5]
- D. [32, 16, 16, 5]

Question 4

*Assume that the tensor before the last tensor of a CNN has shape [32, 32, 32, 10] and we apply **5 filters** each of which has the shape **[5,5,10]** and strides= **[2,2]** with padding = **'same'** to obtain the last tensor. What is the **shape** of the output tensor?*

- ❑ A. [16, 16, 5]
- ❑ B. [14, 14, 5]
- ❑ C. [32, 14, 14, 5]
- ❑ D. [32, 16, 16, 5] **[x]**

- Given an image in a minibatch, we convolve each [32, 32, 10] with [5, 5 ,10] to achieve a 16x16 feature map.

$$\left\lfloor \frac{32-1}{2} \right\rfloor + 1 = 16$$

- There are **5** filters → [16, 16, **5**]
- There are **32** inputs in a minibatch → [**32**, 16, 16, 5]

Question 5

Assume that the tensor before the last tensor of a CNN has shape [32, 32, 32, 10] and we apply ***5 filters*** *each of which has the shape* ***[5,5,10]*** *and strides=* ***[2,2]*** *with padding =* ***'valid'*** *to obtain the last tensor. We flatten this tensor to a* ***fully connected*** *(FC) layer. What is the* ***number of neurons*** *on this FC layer?*

- A. 16 x 16 x5
- B. 14 x 14 x 5
- C. 32 x 16 x 16 x5
- D. 32 x 14 x 14 x 5

Question 5

*Assume that the tensor before the last tensor of a CNN has shape [32, 32, 32, 10] and we apply **5 filters** each of which has the shape **[5,5,10]** and strides= **[2,2]** with padding = **'valid'** to obtain the last tensor. We flatten this tensor to a **fully connected** (FC) layer. What is the **number of neurons** on this FC layer?*

why not 32 x 14 x 14 x 5?

- ❑ A. 16 x 16 x5
- ❑ B. 14 x 14 x 5 **[x]**
- ❑ C. 32 x 16 x 16 x5
- ❑ D. 32 x 14 x 14 x 5

- Given an image in a minibatch, we convolve each [32, 32, 10] with [5, 5 ,10] to achieve a 14x14 feature map.

$$\left\lfloor \frac{32-5}{2} \right\rfloor + 1 = 14$$

- There are **5** filters → [14, 14, **5**]
- We flatten [14, 14, 5] and obtain 14x14x5 neurons

Question 6

What likely happen if using a large filter (e.g., 7x7, 9x9) with a deep model (e.g., 20 layers) if there are few images?

- ❑ A. Overfitting
- ❑ B. Underfitting

Question 6

What likely happen if using a large filter (e.g., 7x7, 9x9) with a deep model (e.g., 20 layers) if there are few images?

- ❑ A. Overfitting **[x]**
- ❑ B. Underfitting

- Larger filter → More parameters → Overfitting problem
- Small filter → Fewer parameters → Underfitting problem
- If we use large enough model (several layers), 3x3 filter will be a common choice.

Question 7

Which is a good CNN model architecture?

- A. Input layer → Convolutional layer (Activation) → Pooling layer → FC layer → Output
- B. Input layer → Pooling layer → Convolutional layer (Activation) → FC layer → Output
- C. Input layer → FC Layer → Pooling layer → Convolutional layer (Activation) → Output
- D. Input layer → Convolutional layer (Activation) → FC layer → Pooling layer → Output

Question 7

Which is a good CNN model architecture?

- ❑ A. Input layer → Convolutional layer (Activation) → Pooling layer → FC layer → Output **[x]**
- ❑ B. Input layer → Pooling layer → Convolutional layer (Activation) → FC layer → Output
- ❑ C. Input layer → FC Layer → Pooling layer → Convolutional layer (Activation) → Output
- ❑ D. Input layer → Convolutional layer (Activation) → FC layer → Pooling layer → Output

Question 8

Given an implementation as below. What is the shape of h1?

```
X = Input(shape=(32, 32, 3))
h1 = Conv2D(filters=10, kernel_size=(3, 3), strides=(1, 1), padding='same')(X)
h1 = AveragePooling2D(pool_size=(2, 2), strides=(2, 2))(h1)
h2 = Conv2D(filters=10, kernel_size=(3, 3), strides=(1, 1), padding='same')(h1)
h2 = AveragePooling2D(pool_size=(2, 2), strides=(2, 2))(h2)
h3 = Conv2D(filters=10, kernel_size=(3, 3), strides=(1, 1), padding='same')(h2)
h4 = Flatten()(h3)
p = Dense(10)(h4)
print("h1", h1.shape)
```

- A. (16,16,3)
- B. (16,16,10)
- C. (None,16,16,3)
- D. (None, 16,16,10)

Question 8

Given an implementation as below. What is the shape of h1?

```
X = Input(shape=(32, 32, 3))
h1 = Conv2D(filters=10, kernel_size=(3, 3), strides=(1, 1), padding='same')(X)
h1 = AveragePooling2D(pool_size=(2, 2), strides=(2, 2))(h1)
h2 = Conv2D(filters=10, kernel_size=(3, 3), strides=(1, 1), padding='same')(h1)
h2 = AveragePooling2D(pool_size=(2, 2), strides=(2, 2))(h2)
h3 = Conv2D(filters=10, kernel_size=(3, 3), strides=(1, 1), padding='same')(h2)
h4 = Flatten()(h3)
p = Dense(10)(h4)
print("h1", h1.shape)
```

- A. (16,16,3)
- B. (16,16,10)
- C. (None,16,16,3)
- D. (None, 16,16,10) **[x]**

before pooling: 32 x 32 x 10
after pooling: 16 x 16 x 10

because no batch size?

Question 9

Given an implementation as below. What is the shape of h1?

```
X = Input(shape=(32, 32, 3))
h1 = Conv2D(filters=10, kernel_size=(3, 3), strides=(1, 1), padding='valid')(X)
h1 = AveragePooling2D(pool_size=(2, 2), strides=(2, 2))(h1)
h2 = Conv2D(filters=10, kernel_size=(3, 3), strides=(1, 1), padding='valid')(h1)
h2 = AveragePooling2D(pool_size=(2, 2), strides=(2, 2))(h2)
h3 = Conv2D(filters=10, kernel_size=(3, 3), strides=(1, 1), padding='valid')(h2)
h4 = Flatten()(h3)
p = Dense(10)(h4)
print("h1", h1.shape)
```

- A. (None,16,16,10)
- B. (None,15,15,10)
- C. (None,14,14,10)
- D. (None, 13,13,10)

Question 9

Given an implementation as below. What is the shape of h1?

```
X = Input(shape=(32, 32, 3))
h1 = Conv2D(filters=10, kernel_size=(3, 3), strides=(1, 1), padding='valid')(X)
h1 = AveragePooling2D(pool_size=(2, 2), strides=(2, 2))(h1)
h2 = Conv2D(filters=10, kernel_size=(3, 3), strides=(1, 1), padding='valid')(h1)
h2 = AveragePooling2D(pool_size=(2, 2), strides=(2, 2))(h2)
h3 = Conv2D(filters=10, kernel_size=(3, 3), strides=(1, 1), padding='valid')(h2)
h4 = Flatten()(h3)
p = Dense(10)(h4)
print("h1", h1.shape)
```

- A. (None,16,16,10)
- B. (None,15,15,10) **[x]**
- C. (None,14,14,10)
- D. (None, 13,13,10)

before pooling: 30x30x10 ; floor(32-3)+1 = 30
after pooling : 15 x 15 x 10

Question 10

Given an implementation as below. What are the shape of h1/h2/h3?

```
X = Input(shape=(32, 32, 3))
h1 = Conv2D(filters=10, kernel_size=(3, 3), strides=(1, 1), padding='valid')(X)
h1 = AveragePooling2D(pool_size=(2, 2), strides=(2, 2))(h1)
h2 = Conv2D(filters=10, kernel_size=(3, 3), strides=(1, 1), padding='valid')(h1)
h2 = AveragePooling2D(pool_size=(2, 2), strides=(2, 2))(h2)
h3 = Conv2D(filters=10, kernel_size=(3, 3), strides=(1, 1), padding='valid')(h2)
h4 = Flatten()(h3)
p = Dense(10)(h4)
print("h1", h1.shape)
print("h2", h2.shape)
print("h3", h3.shape)
```

- A. (None,15,15,10) / (None,6,6,10) / (None,3,3,10)
- B. (None,15,15,10) / (None,6,6,10) / (None,4,4,10)
- C. (None,15,15,10) / (None,7,7,10) / (None,3,3,10)
- D. (None,15,15,10) / (None,7,7,10) / (None,4,4,10)

Question 10

Given an implementation as below. What are the shape of h1/h2/h3?

```
X = Input(shape=(32, 32, 3))
h1 = Conv2D(filters=10, kernel_size=(3, 3), strides=(1, 1), padding='valid')(X)
h1 = AveragePooling2D(pool_size=(2, 2), strides=(2, 2))(h1)
h2 = Conv2D(filters=10, kernel_size=(3, 3), strides=(1, 1), padding='valid')(h1)
h2 = AveragePooling2D(pool_size=(2, 2), strides=(2, 2))(h2)
h3 = Conv2D(filters=10, kernel_size=(3, 3), strides=(1, 1), padding='valid')(h2)
h4 = Flatten()(h3)
p = Dense(10)(h4)
print("h1", h1.shape)
print("h2", h2.shape)
print("h3", h3.shape)
```

- A. (None,15,15,10) / (None,6,6,10) / (None,3,3,10)
- B. (None,15,15,10) / (None,6,6,10) / (None,4,4,10) **[x]**
- C. (None,15,15,10) / (None,7,7,10) / (None,3,3,10)
- D. (None,15,15,10) / (None,7,7,10) / (None,4,4,10)

h1: 32 x 32 x 10
h1 (after pool) : 15 x 15 x10
h2: 12 x 12 x 10
h2: 6 x 6 10
h3: 4 x 4 x 10; floor(6 - 3) / 1 + 1

Question 11

Given an implementation as below. What is the shape of h4?

```
X = Input(shape=(32, 32, 3))
h1 = Conv2D(filters=10, kernel_size=(3, 3), strides=(1, 1), padding='valid')(X)
h1 = AveragePooling2D(pool_size=(2, 2), strides=(2, 2))(h1)
h2 = Conv2D(filters=10, kernel_size=(3, 3), strides=(1, 1), padding='valid')(h1)
h2 = AveragePooling2D(pool_size=(2, 2), strides=(2, 2))(h2)
h3 = Conv2D(filters=10, kernel_size=(3, 3), strides=(1, 1), padding='valid')(h2)
h4 = Flatten()(h3)
p = Dense(10)(h4)
print("h4", h4.shape)
```

- A. (None,4,4,10)
- B. (None,90)
- C. (None,160)
- D. (90,)

Question 11

Given an implementation as below. What is the shape of h4?

```
X = Input(shape=(32, 32, 3))
h1 = Conv2D(filters=10, kernel_size=(3, 3), strides=(1, 1), padding='valid')(X)
h1 = AveragePooling2D(pool_size=(2, 2), strides=(2, 2))(h1)
h2 = Conv2D(filters=10, kernel_size=(3, 3), strides=(1, 1), padding='valid')(h1)
h2 = AveragePooling2D(pool_size=(2, 2), strides=(2, 2))(h2)
h3 = Conv2D(filters=10, kernel_size=(3, 3), strides=(1, 1), padding='valid')(h2)
h4 = Flatten()(h3)
p = Dense(10)(h4)
print("h4", h4.shape)
```

- A. (None,4,4,10)
- B. (None,90)
- C. (None,160) **[x]**
- D. (90,)

Question 12

Given an implementation as below. What is the shape of W1 and b1?

```
X = Input(shape=(32, 32, 3))
h1 = Conv2D(filters=10, kernel_size=(3, 3), strides=(1, 1), padding='valid')(X) # Layer 1
h1 = AveragePooling2D(pool_size=(2, 2), strides=(2, 2))(h1)
h2 = Conv2D(filters=10, kernel_size=(3, 3), strides=(1, 1), padding='valid')(h1)
h2 = AveragePooling2D(pool_size=(2, 2), strides=(2, 2))(h2)
h3 = Conv2D(filters=10, kernel_size=(3, 3), strides=(1, 1), padding='valid')(h2)
h4 = Flatten()(h3)
p = Dense(10)(h4)

model = tf.keras.Model(inputs=X, outputs=p)

W1, b1 = model.layers[1].weights
print(W1.shape)
print(b1.shape)
```

- A. (3,3,3,10), (10,)
- B. (3,3,3,10), (3,3,10)
- C. (15,15,3,10), (10,)
- D. (15,15,3,10), (3,3,10)

Question 12

Given an implementation as below. What is the shape of W1 and b1?

```
X = Input(shape=(32, 32, 3))
h1 = Conv2D(filters=10, kernel_size=(3, 3), strides=(1, 1), padding='valid')(X) # Layer 1
h1 = AveragePooling2D(pool_size=(2, 2), strides=(2, 2))(h1)
h2 = Conv2D(filters=10, kernel_size=(3, 3), strides=(1, 1), padding='valid')(h1)
h2 = AveragePooling2D(pool_size=(2, 2), strides=(2, 2))(h2)
h3 = Conv2D(filters=10, kernel_size=(3, 3), strides=(1, 1), padding='valid')(h2)
h4 = Flatten()(h3)
p = Dense(10)(h4)

model = tf.keras.Model(inputs=X, outputs=p)

W1, b1 = model.layers[1].weights
print(W1.shape)
print(b1.shape)
```

- A. (3,3,3,10), (10,) **[x]**
- B. (3,3,3,10), (3,3,10)
- C. (15,15,3,10), (10,)
- D. (15,15,3,10), (3,3,10)

Weight's shape: (*height*, *width*, *depth*, *num_filters)*
Bias's shape: (*num_filters,)*