

COPYRIGHT WARNING: Copyright in these original lectures is owned by Monash University. You may transcribe, take notes, download or stream lectures for the purpose of your research and study only. If used for any other purpose, (excluding exceptions in the Copyright Act 1969 (Cth)) the University may take legal action for infringement of copyright.

Do not share, redistribute, or upload the lecture to a third party without a written permission!

FIT3181 Deep Learning

Week 09: Representation Learning (II): Autoencoder and Variational Auto-Encoder

Lecturer: Lim Chern Hong

Email: lim.chernhong@monash.edu

Outline

- Revision of some basic knowledge
- Learning efficient representations
- Auto-Encoder
 - Standard Auto-Encoder, Sparse, Contractive, Denoising Auto-Encoders
- Stochastic Auto-Encoder
 - Variational Auto-Encoder
- Further reading recommendation
 - [Hands-On, ch15]
 - [Deep learning, ch14]

Revision of some basic knowledge

Revision of basic knowledge

- Given two **discrete distributions** $p = [p_i]_{i=1}^d$ ($p_i \geq 0$ and $\sum_{i=1}^d p_i = 1$) and $q = [q_i]_{i=1}^d$ ($q_i \geq 0$ and $\sum_{i=1}^d q_i = 1$).

- **Kullback-Leibler (KL) divergence between p, q**

$$KL(p, q) = \sum_{i=1}^d p_i \log \frac{p_i}{q_i}$$

- **Cross-entropy (CE) divergence between p, q**

$CE(p, q) = -\sum_{i=1}^d p_i \log q_i = KL(p, q) + H(p)$ where $H(p) = -\sum_{i=1}^d p_i \log p_i$ is the entropy of p .

- **Cross-entropy between two Bernoulli distributions**

- Given $0 \leq a, b \leq 1$, we have two Bernoulli distributions $Ber(a)$ and $Ber(b)$, the **CE divergence** between them is
$$CE([a, 1 - a], [b, 1 - b]) = -a \log b - (1 - a) \log(1 - b)$$

Revision of basic knowledge

- Given **two continuous distributions** with the **probability density functions** (pdf) $p(x)$ and $q(x)$ respectively
- **KL divergence between p and q**

$$KL(p, q) = \int p(x) \log \frac{p(x)}{q(x)} dx$$


- **Multivariate Gaussian distribution in \mathbb{R}^d**

- $N(x | \mu, \Sigma) = \frac{1}{\det(2\pi\Sigma)^{1/2}} \exp\{-\frac{1}{2}(x - \mu)^T \Sigma^{-1}(x - \mu)\}$

- **KL divergence between two Gaussian distributions in \mathbb{R}^d**

- $KL\left(N\left(\mu, \text{diag}(\sigma^2)\right), N(0, I)\right) = \frac{1}{2}(\|\sigma\|_2^2 + \|\mu\|_2^2 - d - \sum_{i=1}^d \log(\sigma_i^2))$

The derivation of the general case: <https://mr-easy.github.io/2020-04-16-kl-divergence-between-2-gaussian-distributions/>





Learning efficient representation

The importance of efficient/appropriate representation

Roman Numeral Table					
1	I	14	XIV	27	XXVII
2	II	15	XV	28	XXVIII
3	III	16	XVI	29	XXIX
4	IV	17	XVII	30	XXX
5	V	18	XVIII	31	XXXI
6	VI	19	XIX	40	XL
7	VII	20	XX	50	L
8	VIII	21	XXI	60	LX
9	IX	22	XXII	70	LXX
10	X	23	XXIII	80	LXXX
11	XI	24	XXIV	90	XC
12	XII	25	XXV	100	C
13	XIII	26	XXVI	101	CI
				150	CL
				200	CC
				300	CCC
				400	CD
				500	D
				600	DC
				700	DCC
				800	DCCC
				900	CM
				1000	M
				1600	MDC
				1700	MDCC
				1900	MCM

What is LX divided by V?

The importance of efficient/appropriate representation


Roman Numeral Table					
1	I	14	XIV	27	XXVII
2	II	15	XV	28	XXVIII
3	III	16	XVI	29	XXIX
4	IV	17	XVII	30	XXX
5	V	18	XVIII	31	XXXI
6	VI	19	XIX	40	XL
7	VII	20	XX	50	L
8	VIII	21	XXI	60	LX
9	IX	22	XXII	70	LXX
10	X	23	XXIII	80	LXXX
11	XI	24	XXIV	90	XC
12	XII	25	XXV	100	C
13	XIII	26	XXVI	101	CI
				150	CL
				200	CC
				300	CCC
				400	CD
				500	D
				600	DC
				700	DCC
				800	DCCC
				900	CM
				1000	M
				1600	MDC
				1700	MDCC
				1900	MCM

What is 60 divided by 5?

Efficient/appropriate representations



William Chase



Herbert Simon

- The relationship between memory, perception, and pattern matching was studied by William Chase and Herbert Simon in 1970s
 - How can our brain memorise complicated things?
 - How can our brain work out efficient internal representations?

Efficient/appropriate representations

- Expert chess players can memorise the positions of all the pieces in a game within 5 seconds
 - A task most of us would find **impossible**
 - Expert chess players do not have **much better memory** than us. How can they do that?
- This is only the case when all pieces are placed in realistic positions from actual games, not when the pieces are placed in random positions



Efficient/appropriate representations


Expert chess player

Encoding

- Expert chess player do not memorise the positions, they instead memorise the patterns
- Transform all the piece positions to the patterns in memory

Decoding

- Reconstruct the piece positions from the patterns in memory



(Source: Hands-On Ch15)

Auto-Encoder

Encoding (Encoder)

- Transform input data to internal representation or (lossy/lossless) summary

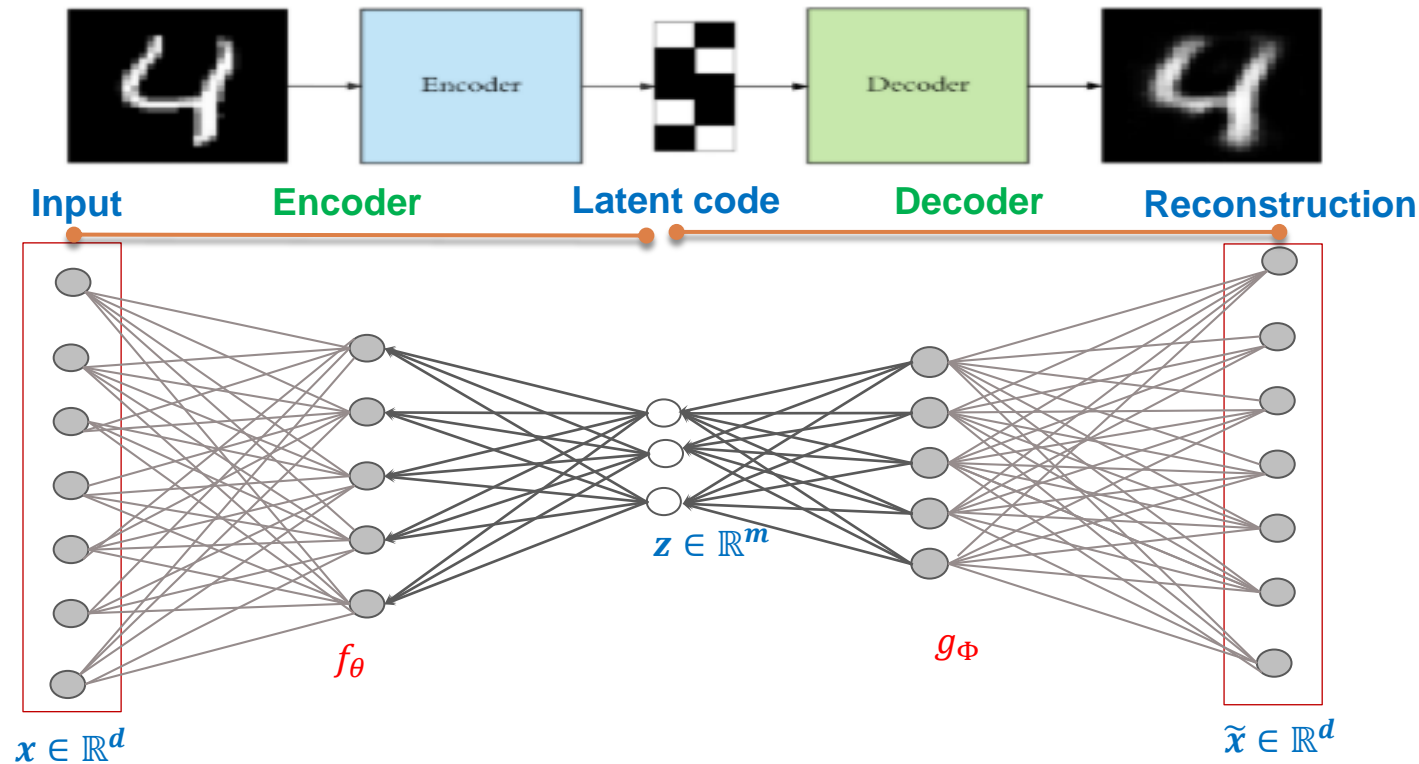
Decoding (Decoder)

- Reconstruct input data from internal representations or (lossy/lossless) summary



Deep Auto-Encoder

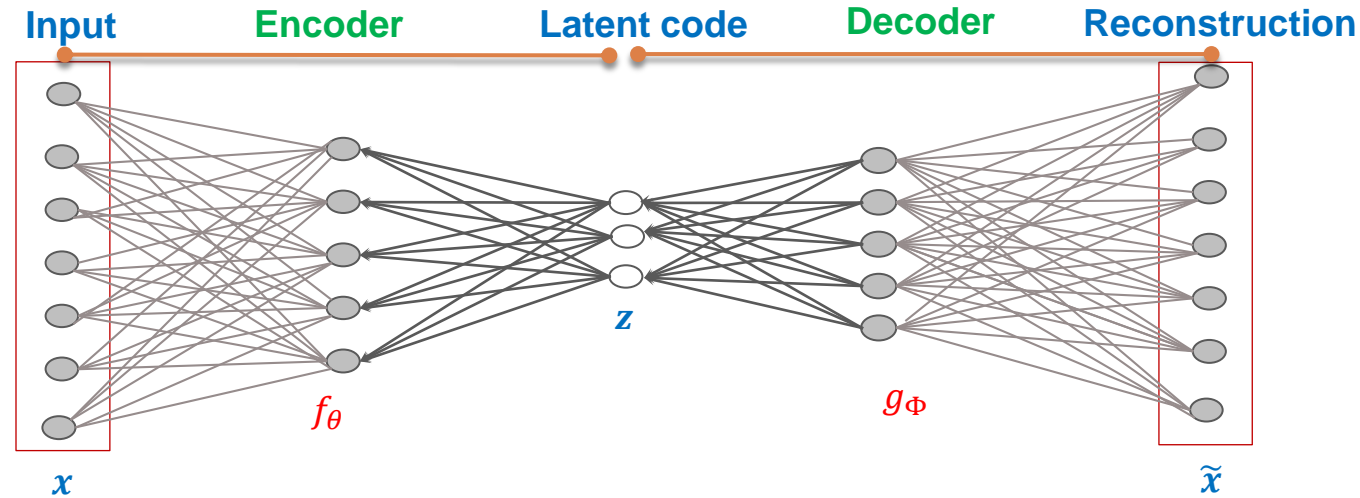
Auto-Encoder



- $x \sim \mathbb{P}$
 - \mathbb{P} is the data distribution over \mathbb{R}^d
- **Encoding**
 - $z = f_\theta(x) \in \mathbb{R}^m$
- **Decoding**
 - $\tilde{x} = g_\phi(z)$ is said to be the **reconstruction** of x .
- How to justify that the **latent code** z can **preserve crucial information** of its input x ?
- How **accurate** we can **reconstruct** x from z ?
 - Reconstruction error: $d(x, \tilde{x})$
 - Distance between x and \tilde{x}

Auto-Encoder


Reconstruction error



- Minimize **reconstruction error** over training set $D = \{x_1, \dots, x_N\}$
 - $\min_{\theta, \Phi} \mathbb{E}_{x \sim \mathbb{P}} [d(x, \tilde{x})] = \min_{\theta, \Phi} \mathbb{E}_{x \sim \mathbb{P}} [d(x, g_\Phi(f_\theta(x)))]$
 - $\min_{\theta, \Phi} \frac{1}{N} \sum_{i=1}^N d(x_i, \tilde{x}_i) = \min_{\theta, \Phi} \frac{1}{N} \sum_{i=1}^N d(x_i, g_\Phi(z_i)) = \min_{\theta, \Phi} \frac{1}{N} \sum_{i=1}^N d(x_i, \underbrace{g_\Phi(\underbrace{f_\theta(x_i)}_{z_i})}_{\tilde{x}_i})$
- How to **define** $d(x, \tilde{x})$?
 - $x, \tilde{x} \in \mathbb{R}^d$: $d(x, \tilde{x}) = \frac{1}{2} \|x - \tilde{x}\|_2^2$ (L2 distance)
 - $x, \tilde{x} \in [0, 1]^d$ (applied **sigmoid** on the output): $d(x, \tilde{x}) = \sum_{i=1}^d CE([x_i, 1 - x_i], [\tilde{x}_i, 1 - \tilde{x}_i])$
$$= \sum_{i=1}^d [-x_i \log \tilde{x}_i - (1 - x_i) \log(1 - \tilde{x}_i)]$$

Undercomplete Auto-Encoder


- Why would one want to learn to copy the input to the output?
 - We are **not actually interested** in the output
 - The hope is that by learning to perform copying from input to output via the intermediate code z , this code will capture useful and key properties of the data
- Ensure the code z learn useful information is through 'compression'
 - Letting dimension of the code z to be smaller than the dimension of the input.
 - This is called an undercomplete AE
- When the decoder is linear and mean squared error loss is used, this is identical to Principal Component Analysis (PCA)



Undercomplete AE when $m < d$

Overcomplete Auto-Encoder

- Overcomplete AE is when the latent code dimension m is greater than the input dimension d .
- In undercomplete AE, $m < d$, hence the code can learn salient features of the data,
 - For overcomplete case, the encoder/decoder could be **too powerful**, hence it can copy (even perfectly) without learning any useful code z !
- How to make overcomplete AE **useful**?
 - Regularization



Overcomplete AE when $m \geq d$

Implementation of Auto-Encoder

```
class GeneralAE:
    def __init__(self, optimizer = keras.optimizers.SGD(lr=0.1)):
        self.encoder = None
        self.decoder = None
        self.auto_encoder = None
        self.optimizer = optimizer

    @staticmethod
    def rounded_accuracy(y_true, y_pred):
        return keras.metrics.binary_accuracy(tf.round(y_true), tf.round(y_pred))

    def encode(self, X=None):
        return self.encoder.predict(X)

    def decode(self, h=None):
        return self.decoder.predict(h)

    def reconstruct(self, X=None):
        self.auto_encoder.predict(X)
```

```
def show_reconstructions(self, X= None, n_cols = 5):
    reconstructions = self.auto_encoder.predict(X)
    n_images = len(X)
    n_rows = math.ceil(n_images/n_cols)
    fig = plt.figure(figsize=(2*n_cols*1.5, n_rows*1.5))
    plt.axis("off")
    for i in range(n_images):
        plt.subplot(n_rows, 2*n_cols, 2*i+1)
        plt.imshow(X[i], cmap="gray")
        plt.xlabel("real")
        plt.xticks([])
        plt.yticks([])
        plt.grid(False)
        plt.subplot(n_rows, 2*n_cols, 2*i+2)
        plt.imshow(reconstructions[i], cmap="gray")
        plt.xlabel("reconstruct")
        plt.xticks([])
        plt.yticks([])
        plt.grid(False)

    def build(self):
        pass

    def train(self, *args, **kwargs):
        self.auto_encoder.fit(*args, **kwargs)
```

Implementation of Standard Auto-Encoder

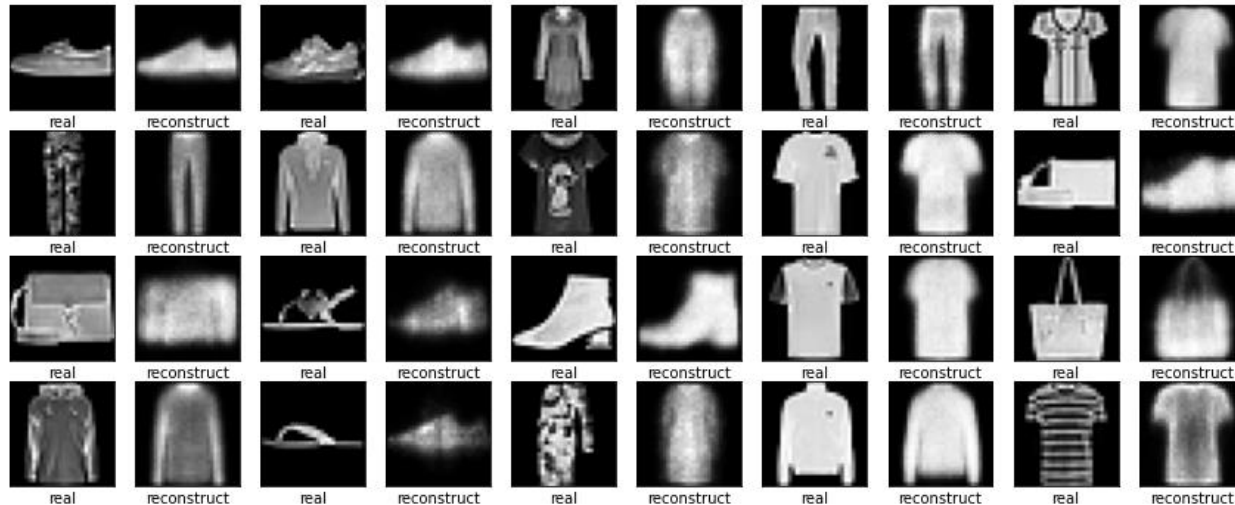
```
class StandardAE(GeneralAE):
    def __init__(self, optimizer = keras.optimizers.SGD(lr=0.1)):
        super(StandardAE, self).__init__(optimizer)

    def build(self):
        self.encoder = keras.models.Sequential([keras.layers.Flatten(input_shape=[28, 28]),
                                                keras.layers.Dense(100, activation="selu"),
                                                keras.layers.Dense(30, activation="selu")])

        self.decoder = keras.models.Sequential([keras.layers.Dense(100, activation="selu", input_shape=[30]),
                                                keras.layers.Dense(28 * 28, activation="sigmoid"),
                                                keras.layers.Reshape([28, 28])])


        self.auto_encoder = keras.models.Sequential([self.encoder, self.decoder])
        self.auto_encoder.compile(loss="binary_crossentropy", optimizer=self.optimizer, metrics=[GeneralAE.rounded_accuracy])
```

```
standard_ae.show_reconstructions(X_valid[:20], n_cols=5)
```



Sparse Auto-Encoder

- Given a latent code \mathbf{z} , the sparsity of \mathbf{z} is defined as the number of **zero elements** in \mathbf{z} or $m - \|\mathbf{z}\|_0$
 - $\mathbf{z}_1 = [1.2, 0, -1, 0, 1] \rightarrow \text{sparsity}(\mathbf{z}_1) = 2$
 - $\mathbf{z}_2 = [2.4, 0, -2, 0, 0] \rightarrow \text{sparsity}(\mathbf{z}_2) = 3$
 - In general, **sparser** \mathbf{z} is, **more elements around 0** it has
- We want to find **sparse representation** of the latent code \mathbf{z} of \mathbf{x} that is still able to reconstruct well \mathbf{x} .
 - Hope that the training would **eliminate redundant** elements
- **Sparse auto-encoder**
 - $\min_{\theta, \Phi} \mathbb{E}_{\mathbf{x} \sim \mathbb{P}} [d(\mathbf{x}, g_{\Phi}(f_{\theta}(\mathbf{x}))) + \lambda \Omega(\mathbf{z})]$
 - $\Omega(\mathbf{z})$ is a regularization which is usually a norm over \mathbf{z}
 - $\lambda > 0$ is regularization parameter.



Sparse Auto-Encoder

Regularization

□ Sparse auto-encoder

- $\min_{\theta, \Phi} \mathbb{E}_{\mathbf{x} \sim \mathbb{P}} [d(\mathbf{x}, g_{\Phi}(f_{\theta}(\mathbf{x})))] + \lambda \Omega(\mathbf{z})$
- $\Omega(\mathbf{z})$ is a regularization which is usually a norm over \mathbf{z} , $\lambda > 0$ is regularization parameter.

□ Possible choices for the regularization $\Omega(\mathbf{z})$

- **Norm-1:** $\Omega(\mathbf{z}) = \|\mathbf{z}\|_1 = \sum_{i=1}^d |z_i|$
- **Norm-2:** $\Omega(\mathbf{z}) = \frac{1}{2} \|\mathbf{z}\|_2^2 = \frac{1}{2} \sum_{i=1}^d z_i^2$
- **CE divergence:** $\Omega(\mathbf{z}) = d^{-1} \sum_{i=1}^d CE([z_i, 1 - z_i], [a, 1 - a]) = -\frac{1}{d} \sum_{i=1}^d [z_i \log a + (1 - z_i) \log(1 - a)]$
 - $0 < a < 1$ is a very small number
 - CE is the **cross-entropy loss**

Implementation of Sparse Auto-Encoder


```
class SparseAE(GeneralAE):
    def __init__(self, optimizer = keras.optimizers.SGD(lr=0.1), regularizer = keras.regularizers.l1(l=0.01)):
        super(SparseAE, self).__init__(optimizer)
        self.regularizer = regularizer

    def build(self):
        self.encoder = keras.models.Sequential([keras.layers.Flatten(input_shape=[28, 28]),
                                                keras.layers.Dense(100, activation="selu"),
                                                keras.layers.Dense(30, activation="selu", activity_regularizer=self.regularizer)])

        self.decoder = keras.models.Sequential([keras.layers.Dense(100, activation="selu", input_shape=[30]),
                                                keras.layers.Dense(28 * 28, activation="sigmoid"),
                                                keras.layers.Reshape([28, 28])])

        self.auto_encoder = keras.models.Sequential([self.encoder, self.decoder])
        self.auto_encoder.compile(loss="binary_crossentropy", optimizer=self.optimizer, metrics=[GeneralAE.rounded_accuracy])
```

```
l1_sparse_ae.show_reconstructions(X_valid[:20], n_cols=5)
```



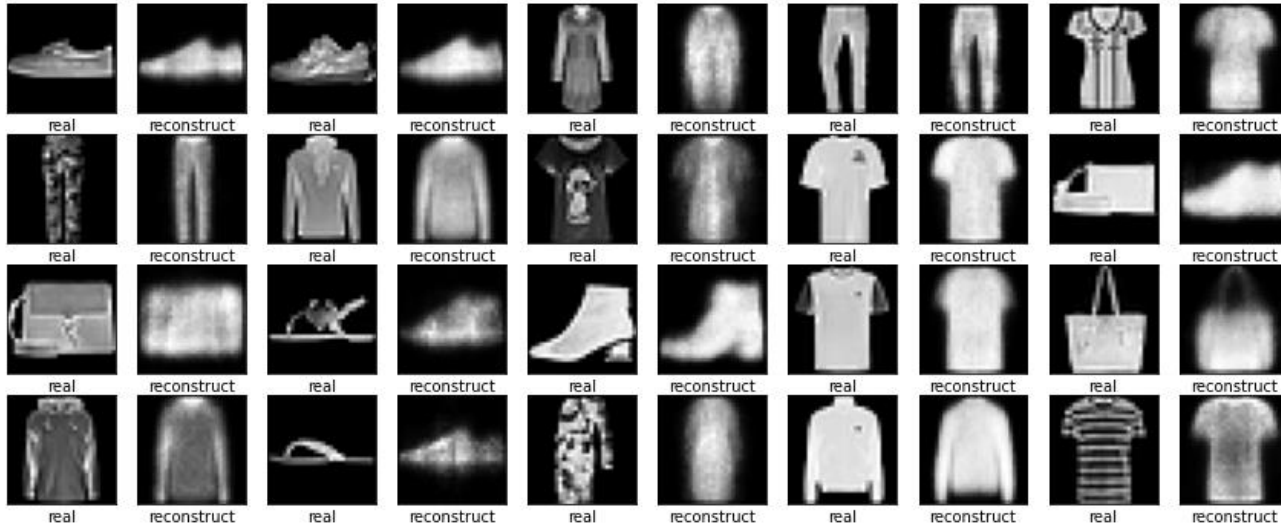
Implementation of Sparse Auto-Encoder

```
kl_divergence = keras.losses.kullback_leibler_divergence

class KLDivergenceRegularizer(keras.regularizers.Regularizer):
    def __init__(self, l, target=0.1):
        self.weight = l
        self.target = target
    def __call__(self, inputs):
        mean_activities = tf.reduce_mean(inputs, axis=0)
        return self.weight * (
            kl_divergence(self.target, mean_activities) + kl_divergence(1. - self.target, 1. - mean_activities))
```

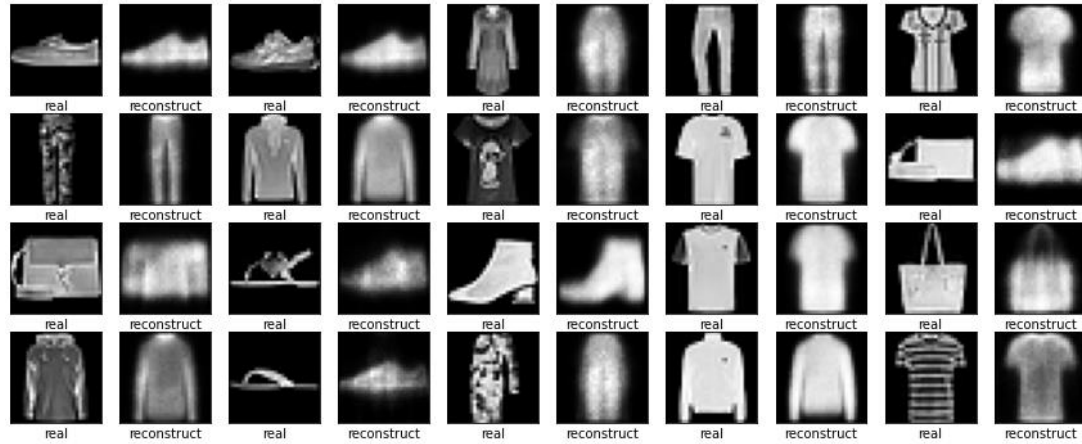
```
kl_sparse_ae = SparseAE(regularizer=KLDivergenceRegularizer(l=0.01))
kl_sparse_ae.build()
kl_sparse_ae.train(X_train, X_train, epochs=20, validation_data=(X_valid, X_valid))
```

```
kl_sparse_ae.show_reconstructions(X_valid[:20], n_cols=5)
```




Sparse Auto-Encoder


Effect of regularization parameter




$\lambda = 0.001$



$\lambda = 0.01$



$\lambda = 0.05$



$\lambda = 10$

Contractive Auto-Encoder


Regularized AE:

- $$\min_{\theta, \Phi} \mathbb{E}_{x \sim \mathbb{P}} [d(x, g_{\Phi}(f_{\theta}(x)))] + \lambda \Omega(x, z)$$

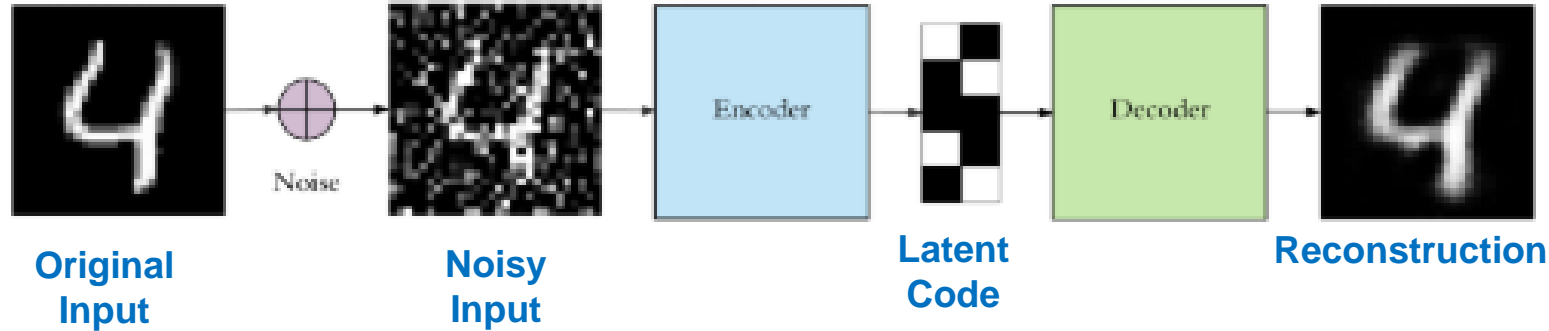
where $\Omega(x, z) = \sum_{i=1}^m \|\nabla_x \mathbf{z}_i\|^2 = \left\| \frac{\partial f_{\theta}(x)}{\partial x} \right\|_F^2$.

- Hence, we train to resist the perturbations of the input by **minimizing the magnitude of the gradient of the encoder f**

- This **contracts** the input neighbourhood to a smaller output neighbourhood, hence the name Contractive AE



Denoising Auto-Encoder



- Add a small **Gaussian noise** to original input and require the auto-encoder to reconstruct the original input
 - $x' = x + \epsilon$ where $\epsilon \sim N(\mathbf{0}, \eta I)$ and learn such that $g_{\Phi}(f_{\theta}(x')) \approx x$
- **Denoising auto-encoder**
 - $\min_{\theta, \Phi} \mathbb{E}_{x \sim \mathbb{P}} [\mathbb{E}_{x' \sim N(x, \eta I)} [d(x, g_{\Phi}(f_{\theta}(x')))]]$

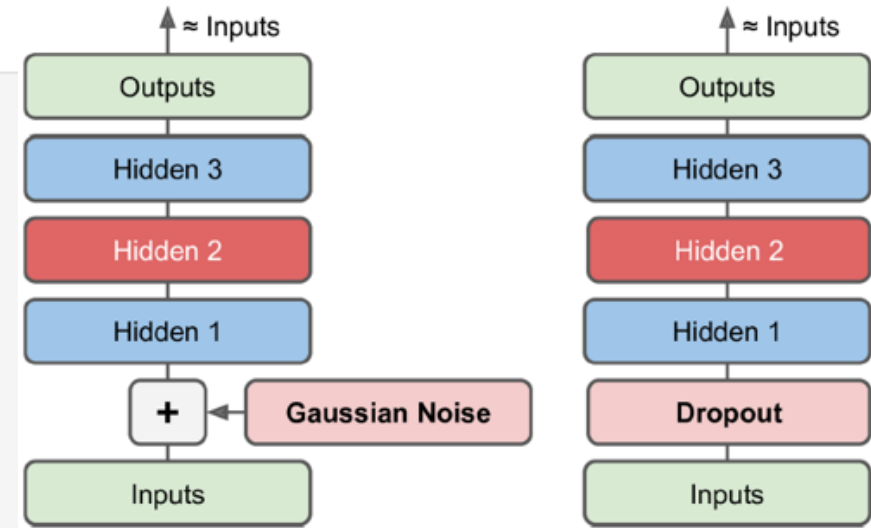
Implementation of Denoising Auto-Encoder

```
class DenoisingAE(GeneralAE):
    def __init__(self, optimizer = keras.optimizers.SGD(lr=0.1), noise= 0.2, seed= 6789):
        super(DenoisingAE, self).__init__(optimizer)
        self.noise = noise
        self.seed = seed
        tf.random.set_seed(self.seed)
        np.random.seed(self.seed)

    def build(self):
        self.encoder = keras.models.Sequential([keras.layers.Flatten(input_shape=[28, 28]),
                                                keras.layers.GaussianNoise(self.noise),
                                                keras.layers.Dense(100, activation="selu"),
                                                keras.layers.Dense(30, activation="selu")])

        self.decoder = keras.models.Sequential([keras.layers.Dense(100, activation="selu", input_shape=[30]),
                                                keras.layers.Dense(28 * 28, activation="sigmoid"),
                                                keras.layers.Reshape([28, 28])])


        self.auto_encoder = keras.models.Sequential([self.encoder, self.decoder])
        self.auto_encoder.compile(loss="binary_crossentropy", optimizer=self.optimizer, metrics=[GeneralAE.rounded_accuracy])
```



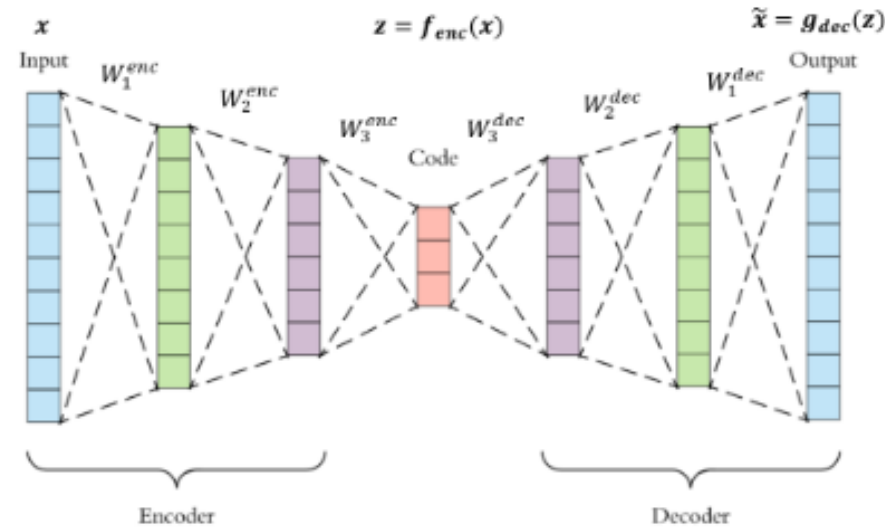
(Source: Hand-On, Ch15)

Implementation of Denoising Auto-Encoder

```
noise = keras.layers.GaussianNoise(0.2)  
denoise_ae.show_reconstructions(noise(X_valid[0:20]), training= True), n_cols=5)
```




Tying encoders and decoders



- To reduce the **number of parameters** and put constraints on Auto-Encoder to avoid trivial models, we can tie encoders and decoders
 - $W_1^{enc} = (W_1^{dec})^T$, $W_2^{enc} = (W_2^{dec})^T$, and $W_3^{enc} = (W_3^{dec})^T$

Representation power, size and depth

- We can build deep AE where both f and g are deep NNs
 - Hence, enjoying the power of deep NNs, especially the universal function approximation properties
 - i.e., even with single layer and one additional hidden layer for the encoder, AE can be very powerful if enough hidden units are given
 - However, without proper regularization, too powerful encoder and decoder are not necessarily good
- Depth can exponentially reduce computational cost and amount of training data in some cases
- Much better compression can be achieved with deep AE compared with shallow or linear AE (Hinton and Salakhutdinov, '06)






Stochastic Auto-Encoder (Variational Auto-Encoder)

Variational Auto-Encoder (VAE)

Generative model viewpoint

- Given a training set $D = \{\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_N\}$ where each $\mathbf{x}_i \sim p_d(\mathbf{x})$.
 - $p_d(\mathbf{x})$ exists but unknown.
- Learn a stochastic decoder $p_\theta(\mathbf{x} | \mathbf{z}) = N(\mathbf{x} | \mu_\theta(\mathbf{z}), \sigma^2 \mathbf{I})$
 - $\mathbf{z} \sim p(\mathbf{z}) = N(0, \mathbf{I}) \rightarrow \tilde{\mathbf{x}} \sim p_\theta(\mathbf{x} | \mathbf{z})$ or $\tilde{\mathbf{x}} = \mu_\theta(\mathbf{z}) + \epsilon \sigma \mathbf{I}$ with $\epsilon \sim N(0, \mathbf{I})$
 - $\tilde{\mathbf{x}}$ seems to **be sampled** from $p_d(\mathbf{x})$.
 - $\tilde{\mathbf{x}}$ mimics true data samples in the training set D .



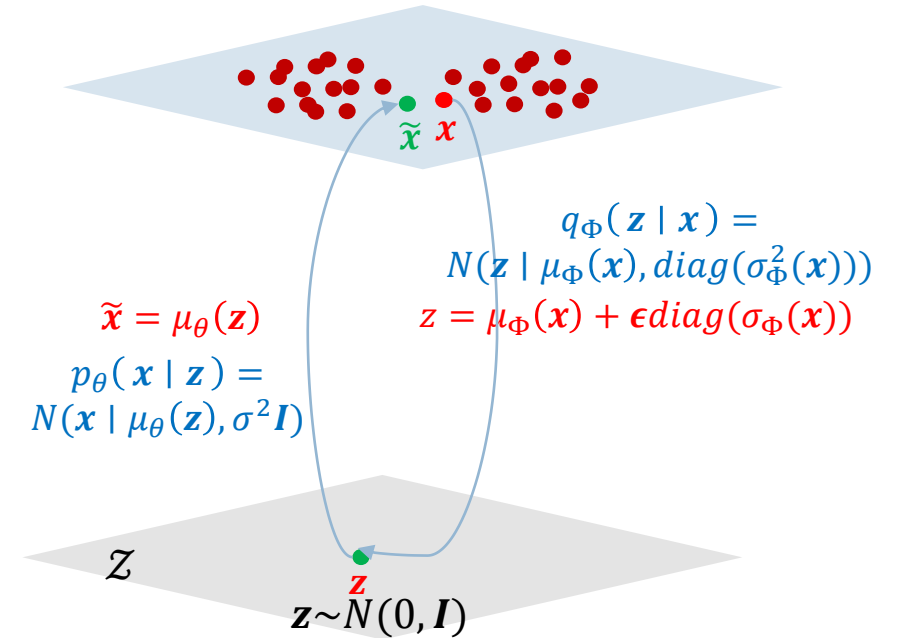

Variational Auto-Encoder (VAE)

□ Use stochastic encoder

- $q_{\Phi}(\mathbf{z} | \mathbf{x}) = N(\mathbf{z} | \mu_{\Phi}(\mathbf{x}), \text{diag}(\sigma_{\Phi}^2(\mathbf{x})))$
- $\mathbf{z} = \mu_{\Phi}(\mathbf{x}) + \epsilon \text{diag}(\sigma_{\Phi}(\mathbf{x}))$ with $\epsilon \sim N(0, \mathbf{I})$

□ Use stochastic decoder

- $p_{\theta}(\mathbf{x} | \mathbf{z}) = N(\mathbf{x} | \mu_{\theta}(\mathbf{z}), \sigma^2 \mathbf{I})$
- $\tilde{\mathbf{x}} = \mu_{\theta}(\mathbf{z})$



Variational Auto-Encoder (VAE)

Objective function

Reconstruction:

$$\max_{\theta, \Phi} \mathbb{E}_x [\log p_\theta(x | z)] \leftrightarrow \min_{\theta, \Phi} \frac{1}{2\sigma^2} \mathbb{E}_x [d(x, \mu_\theta(z))]$$

for $z = \mu_\Phi(x) + \epsilon \text{diag}(\sigma_\Phi(x))^{1/2}$.

Matching to the prior $p(z) = N(0, I)$


$$\min_{\theta, \Phi} \mathbb{E}_x [KL(q_\Phi(z | x) \| N(0, I))] = \frac{1}{2} \min_{\theta, \Phi} \mathbb{E}_x \left[-d + \|\mu_\Phi(x)\|_2^2 + \|\sigma_\Phi(x)\|_2^2 - \text{sum}(\log[\sigma_\Phi(x)^2]) \right]$$

The final objective function

$$\min_{\theta, \Phi} \mathbb{E}_x \left[\mathbb{E}_\epsilon \left[\frac{1}{\sigma^2} d(x, \mu_\theta(z)) + \|\mu_\Phi(x)\|_2^2 + \|\sigma_\Phi(x)\|_2^2 - \text{sum}(\log[\sigma_\Phi(x)^2]) \right] \right] \text{ with } z = \mu_\Phi(x) + \epsilon \text{diag}(\sigma_\Phi(x))^{1/2} \text{ and } \epsilon \sim N(0, I).$$


Use stochastic encoder

- $q_\Phi(z | x) = N(z | \mu_\Phi(x), \text{diag}(\sigma_\Phi^2(x)))$
- $z = \mu_\Phi(x) + \epsilon \text{diag}(\sigma_\Phi(x))$ with $\epsilon \sim N(0, I)$



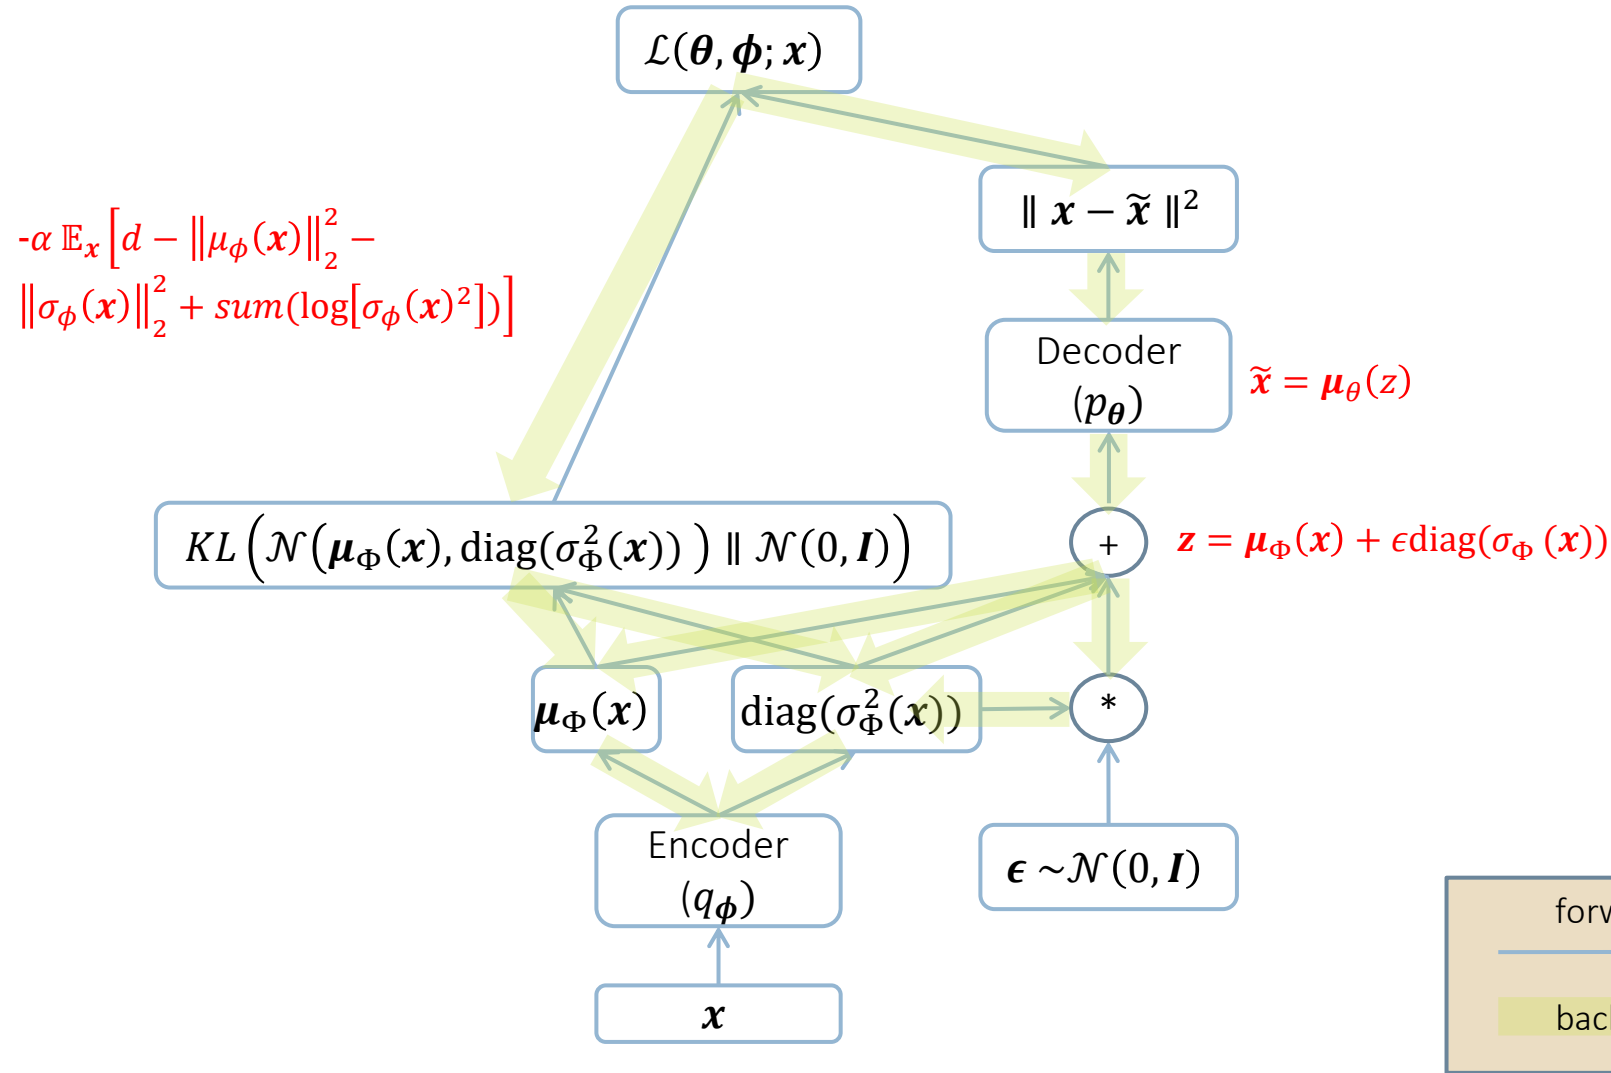
Use stochastic decoder

- $p_\theta(x | z) = N(x | \mu_\theta(z), \sigma^2 I)$
- $\tilde{x} = \mu_\theta(z)$




By matching $q_\Phi(z | x)$ and the prior $N(0, I)$, $z \sim N(0, I)$ seems to be sampled from $q_\Phi(z | x)$ for some x , hence leading to $\tilde{x} = \mu_\theta(z)$ to be **good** and **high-quality** reconstructed examples.

VAE computational graph



VAE Formal Derivation

Not in assessment



Auto-Encoding Variational Bayes

Diederik P. Kingma
Machine Learning Group
Universiteit van Amsterdam
dpkingma@gmail.com

Max Welling
Machine Learning Group
Universiteit van Amsterdam
welling.max@gmail.com

Paper: <https://arxiv.org/pdf/1312.6114.pdf>

- $\log p_\theta(x) = \log \frac{p_\theta(x,z)}{p_\theta(z|x)} = \log \frac{p_\theta(x,z)}{q_\Phi(z|x)} \frac{q_\Phi(z|x)}{p_\theta(z|x)} = \log \frac{p_\theta(x,z)}{q_\Phi(z|x)} + \log \frac{q_\Phi(z|x)}{p_\theta(z|x)}$
- $\mathbb{E}_{x \sim p_d(x)} [\log p_\theta(x)] = \mathbb{E}_{x \sim p_d(x), z \sim q_\Phi(z|x)} [\log p_\theta(x)] = \mathbb{E}_{x \sim p_d(x), z \sim q_\Phi(z|x)} \left[\log \frac{p_\theta(x,z)}{q_\Phi(z|x)} \right] + \mathbb{E}_{x \sim p_d(x), z \sim q_\Phi(z|x)} \left[\log \frac{q_\Phi(z|x)}{p_\theta(z|x)} \right] =$
 $\mathbb{E}_{x \sim p_d(x), z \sim q_\Phi(z|x)} \left[\log \frac{p_\theta(x,z)}{q_\Phi(z|x)} \right] + \mathbb{E}_x \left[\underbrace{\mathbb{E}_{z \sim q_\Phi(z|x)} \left[\log \frac{q_\Phi(z|x)}{p_\theta(z|x)} \right]}_{KL(q_\Phi(z|x) \| p_\theta(z|x)) \geq 0} \right] \geq \mathbb{E}_{x \sim p_d(x), z \sim q_\Phi(z|x)} \left[\log \frac{p_\theta(x,z)}{q_\Phi(z|x)} \right].$
- $\mathbb{E}_{x \sim p_d(x)} [\log p_\theta(x)] \geq \mathbb{E}_{x \sim p_d(x), z \sim q_\Phi(z|x)} \left[\log \frac{p_\theta(x,z)}{q_\Phi(z|x)} \right] = \mathbb{E}_x \left[\mathbb{E}_z \left[\log \frac{p_\theta(x,z)}{q_\Phi(z|x)} \right] \right] = \mathbb{E}_x \left[\mathbb{E}_z [\log p_\theta(x|z)] \right] + \mathbb{E}_x \left[\mathbb{E}_z \left[\log \frac{p(z)}{q_\Phi(z|x)} \right] \right] =$
 $\mathbb{E}_x \left[\mathbb{E}_z [\log p_\theta(x|z)] \right] - \mathbb{E}_x [KL(q_\Phi(z|x) \| N(\mathbf{0}, I))] = \mathbb{E}_x \left[\mathbb{E}_{\epsilon \sim N(\mathbf{0}, I)} [\log p_\theta(x|z)] \right] - \mathbb{E}_x [KL(q_\Phi(z|x) \| N(\mathbf{0}, I))].$
 - Use reparameterization trick: $\mathbb{E}_x [\mathbb{E}_z [\log p_\theta(x|z)]] = \mathbb{E}_x \left[\mathbb{E}_{\epsilon \sim N(\mathbf{0}, I)} [\log p_\theta(x|z)] \right]$ where $z = \mu_\Phi(x) + \epsilon \text{diag}(\sigma_\Phi(x))^{\frac{1}{2}}$.
- The final objective function

$$\max_{\theta, \Phi} \left\{ \mathbb{E}_x \left[\mathbb{E}_{\epsilon \sim N(\mathbf{0}, I)} [\log p_\theta(x|z)] \right] - \mathbb{E}_x [KL(q_\Phi(z|x) \| N(\mathbf{0}, I))] \right\}$$

$$\max_{\theta, \Phi} \left\{ \mathbb{E}_x \left[\mathbb{E}_{\epsilon \sim N(\mathbf{0}, I)} \left[\frac{-\|x - \mu_\Phi(z)\|^2}{2\sigma^2} \right] \right] - \mathbb{E}_x [KL(q_\Phi(z|x) \| N(\mathbf{0}, I))] \right\}$$

Implementation of VAE

```
class VariationalAE(GeneralAE):
    def __init__(self, optimizer = keras.optimizers.SGD(lr=0.1), alpha= 0.5, seed= 6789):
        super(VariationalAE, self).__init__(optimizer)
        self.alpha = alpha
        self.seed = seed
        tf.random.set_seed(self.seed)
        np.random.seed(self.seed)

    def build(self):
        codings_size = 10
        inputs = keras.layers.Input(shape=[28, 28])
        z = keras.layers.Flatten()(inputs)
        z = keras.layers.Dense(150, activation="selu")(z)
        z = keras.layers.Dense(100, activation="selu")(z)
        codings_mean = keras.layers.Dense(codings_size)(z)
        codings_log_var = keras.layers.Dense(codings_size)(z)
        codings = Sampling()(codings_mean, codings_log_var)
        self.encoder = keras.models.Model(inputs=[inputs], outputs=[codings_mean, codings_log_var, codings])

        decoder_inputs = keras.layers.Input(shape=[codings_size])
        x = keras.layers.Dense(100, activation="selu")(decoder_inputs)
        x = keras.layers.Dense(150, activation="selu")(x)
        x = keras.layers.Dense(28 * 28, activation="sigmoid")(x)
        outputs = keras.layers.Reshape([28, 28])(x)
        self.decoder = keras.models.Model(inputs=[decoder_inputs], outputs=[outputs])

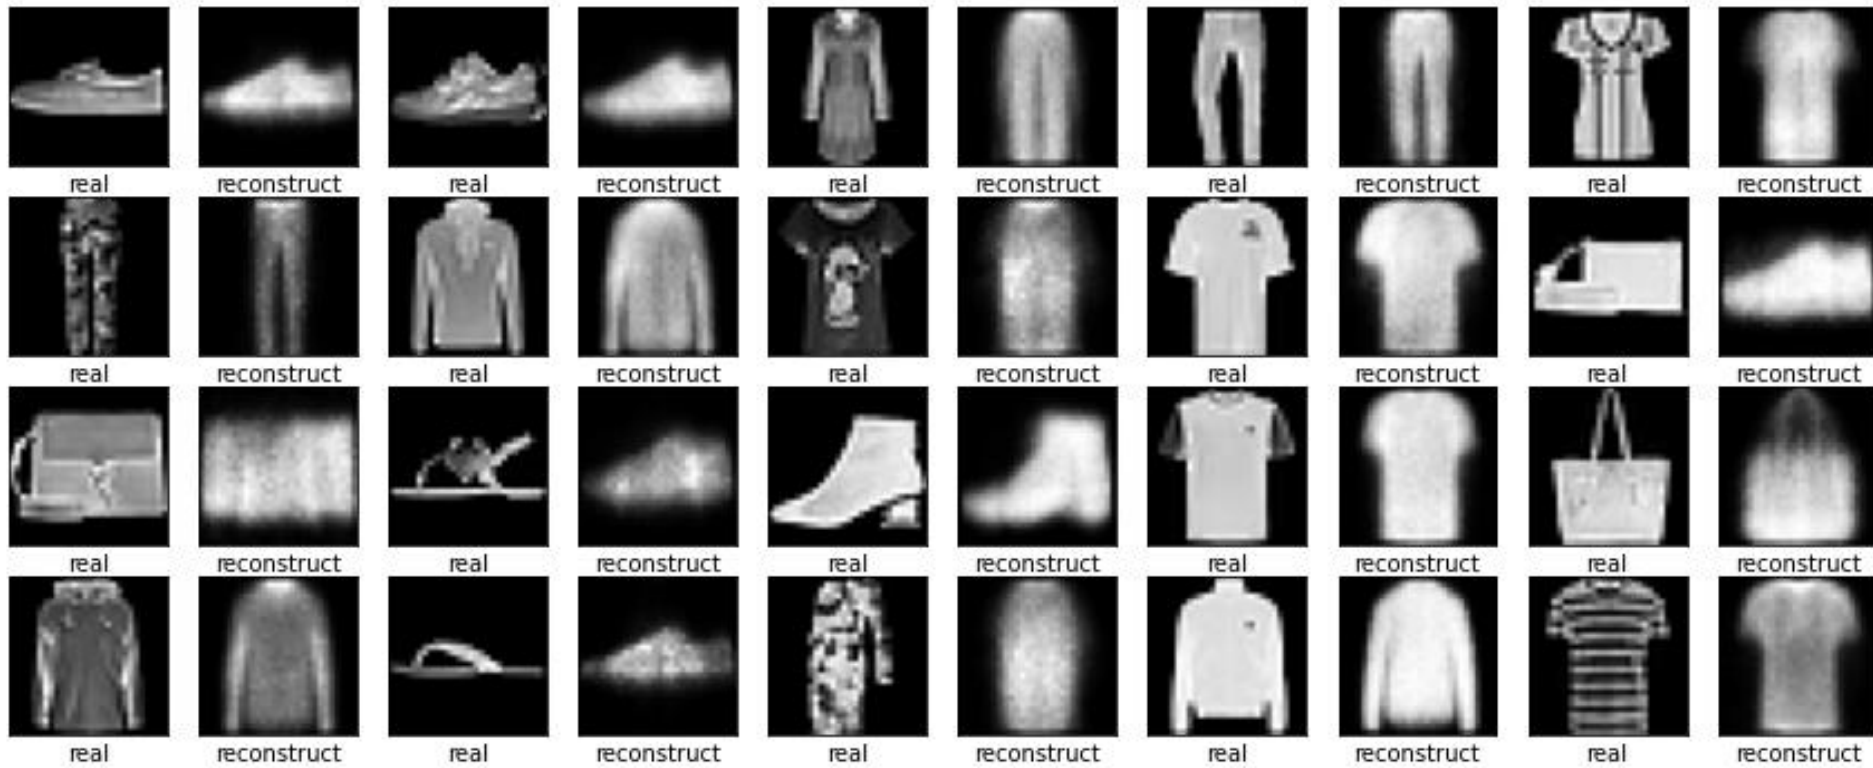
        _, _, codings = self.encoder(inputs)
        reconstructions = self.decoder(codings)
        self.auto_encoder = keras.models.Model(inputs=[inputs], outputs=[reconstructions])

        latent_loss = -self.alpha * tf.reduce_sum(1 + codings_log_var - tf.math.exp(codings_log_var) - tf.math.square(codings_mean), axis=-1)
        self.auto_encoder.add_loss(tf.math.reduce_mean(latent_loss) / 784.)
        self.auto_encoder.compile(loss="binary_crossentropy", optimizer=self.optimizer, metrics=[GeneralAE.rounded_accuracy])
```

```
class Sampling(keras.layers.Layer):
    def call(self, inputs):
        mean, log_var = inputs
        return tf.random.normal(tf.shape(log_var)) * tf.math.exp(log_var / 2) + mean
```

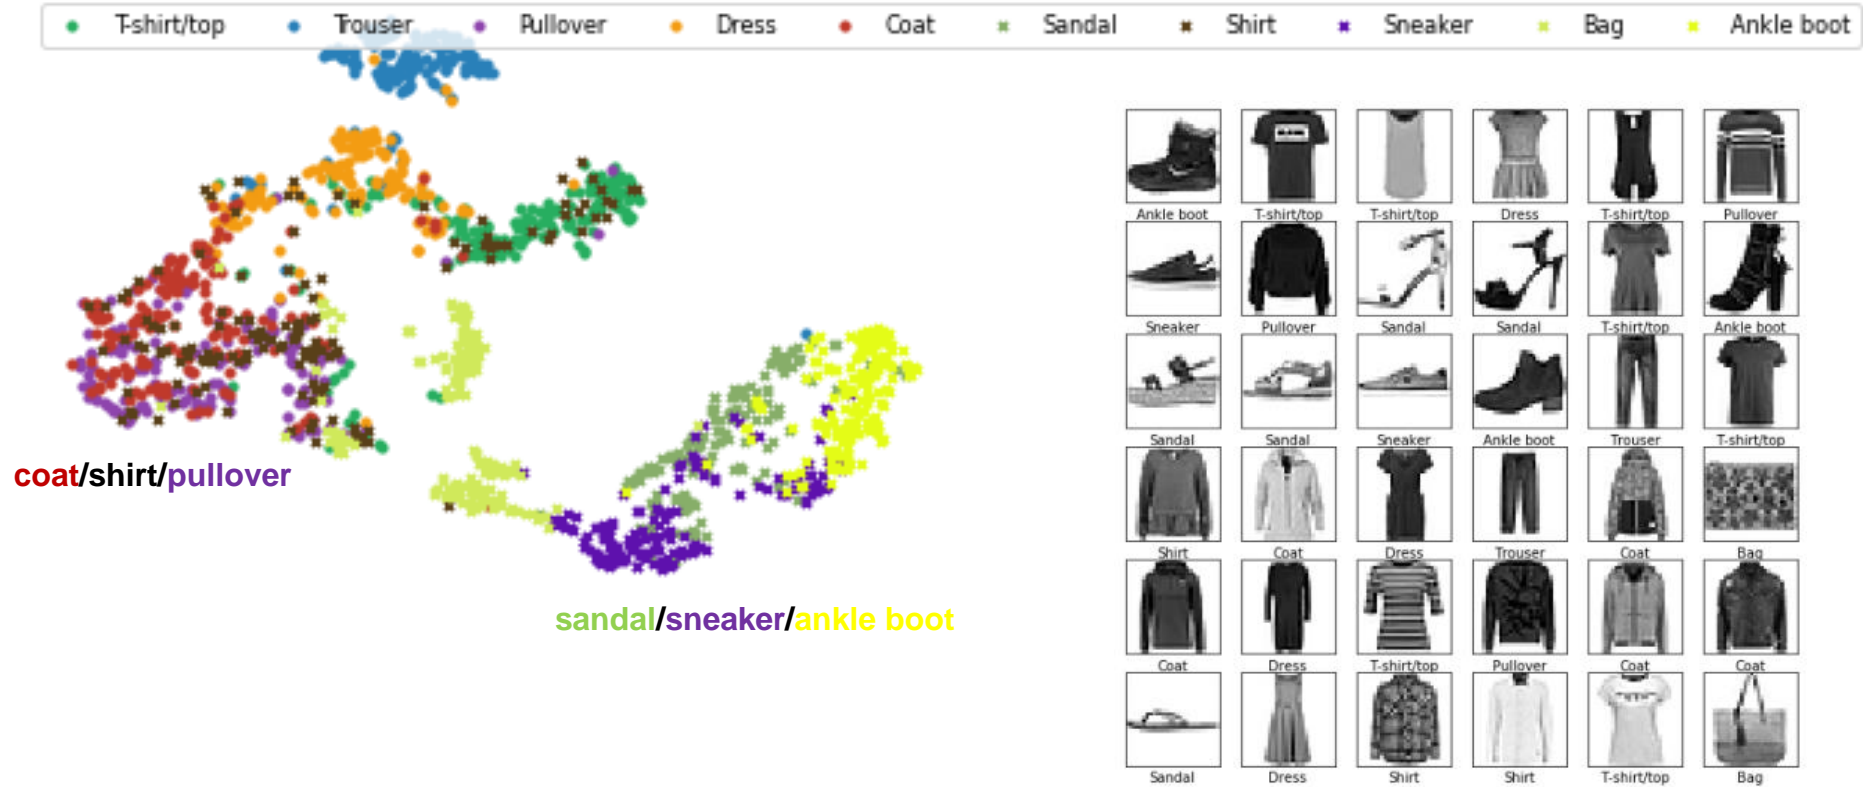
Implementation of VAE

```
vae.show_reconstructions(X_valid[:20], n_cols=5)
```




VAE latent codes with TSNE

```
visualize_latent_codes(X_valid[0:1500], y_valid[0:1500], vae)
```




Variational Auto-encoders (VAE)

- Run VAE on MNIST data with 2-dimensional latent variable \mathbf{z}




sampling

$$p(x | z)$$



$$q(z | x)$$



sampling on manifold

Variational Auto-encoders (VAE)

- VAE weakness: tend to generate blurry images on complex images
- Why? It captures global structures well, but has difficulty with modelling the local information – GAN will be able to address this!



Summary

- Revision of some basic knowledge
- Learning efficient representations
- Auto-Encoder
 - Standard Auto-Encoder, Sparse, Contractive, Denoising Auto-Encoders
- Stochastic Auto-Encoder
 - Variational Auto-Encoder

Thanks for your attention!
Question time

